

VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum -590014, Karnataka.



LAB REPORT on

BIG DATA ANALYTICS

Submitted by

Ishan Bhandari (1BM19CS198)

in partial fulfillment for the award of the degree of
BACHELOR OF ENGINEERING
in
COMPUTER SCIENCE AND ENGINEERING



B.M.S. COLLEGE OF ENGINEERING
(Autonomous Institution under VTU)
BENGALURU-560019
May-2022 to July-2022

B. M. S. College of Engineering,
Bull Temple Road, Bangalore 560019
(Affiliated To Visvesvaraya Technological University, Belgaum)
Department of Computer Science and Engineering



CERTIFICATE

This is to certify that the Lab work entitled “LAB COURSE **BIG DATA ANALYTICS** ” **carried** out by **Ishan Bhandari (IBM19CS198)**, who is a bonafide student of **B. M. S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum during the year 2022. The Lab report has been approved as it satisfies the academic requirements in respect of a **Big Data Analytics - (20CS6PEBDA)** work prescribed for the said degree.

Dr. Pallavi G.B
Assistant Professor
Department of CSE
BMSCE, Bengaluru

Dr. Jyothi S Nayak
Professor and Head
Department of CSE
BMSCE, Bengaluru

Index Sheet

Sl. No.	Experiment Title	Page No.
1	Perform Employee DB operations using Cassandra	4
2	Perform Library DB operations using Cassandra.	9
3	MongoDB- CRUD Demonstration	12
4	Screenshot of Hadoop installed	21
5	Execution of HDFS Commands for interaction with Hadoop Environment.	22
6	Create a Map Reduce program to a) find average temperature for each year from the NCDC data set. b) find the mean max temperature for every month	31
7	For a given Text file, Create a Map Reduce program to sort the content in an alphabetic order	38
8	Create a Map Reduce program to demonstrating join operation	42
9	Program to print word count on scala shell and print “Hello world” on scala IDE	51
10	Using RDD and FlamaMap count how many times each word appears in a file and write out a list of words whose count is strictly greater than 4 using Spark	53

Course Outcome

CO1	Apply the concept of NoSQL, Hadoop or Spark for a given task
CO2	Analyze the Big Data and obtain insight using data analytics mechanisms.
CO3	Design and implement Big data applications by applying NoSQL, Hadoop or Spark

PROGRAM 1

Perform the following DB operations using Cassandra.

1. Create a keyspace by name Employee

2. Create a column family by name

Employee-Info with attributes

Emp_Id Primary Key, Emp_Name,

Designation, Date_of_Joining, Salary, Dept_Name

3. Insert the values into the table in batch

4. Update Employee name and Department of Emp-Id 121

5. Sort the details of Employee records based on salary

6. Alter the schema of the table Employee_Info to add a column Projects which stores a set of Projects done by the corresponding Employee.

7. Update the altered table to add project names.

8. Create a TTL of 15 seconds to display the values of Employees.

```
cqlsh:employee> CREATE KEYSPACE employee WITH REPLICATION={ 'class' :  
'SimpleStrategy',
```

```
'replication_factor' : 1 };
```

```
cqlsh:employee> USE employee;
```

```
cqlsh:employee> create table employee_info(emp_id int PRIMARY KEY, emp_name text,  
... designation text, date_of_joining timestamp, salary double PRIMARY KEY, dept_name text);
```

```
cqlsh:employee> CREATE TABLE employee_info(emp_id int, emp_name text, designation text,  
date_of_joining
```

```
timestamp, salary double, dept_name text, PRIMARY KEY(emp_id, salary));
```

```
cqlsh:employee> BEGIN BATCH INSERT INTO
```

```
... employee_info(emp_id,emp_name,designation,date_of_joining,salary,dept_name)
```

```
... VALUES(100,'John','MANAGER','2021-09-11',30000,'TESTING');
```

```
... INSERT INTO
```

```
... employee_info(emp_id,emp_name,designation,date_of_joining,salary,dept_name)
```

```
... VALUES(111,'Tom','ASSOCIATE','2021-06-22',25000,'DEVELOPING');
```

```

... INSERT INTO
... employee_info(emp_id,emp_name,designation,date_of_joining,salary,dept_name)
... VALUES(121,'Elsa','MANAGER','2021-03-30',35000,'HR');
... INSERT INTO
... employee_info(emp_id,emp_name,designation,date_of_joining,salary,dept_name)
... VALUES(115,'Chris','ASSISTANT','2021-12-30',20000,'DEVELOPING');
... INSERT INTO
... employee_info(emp_id,emp_name,designation,date_of_joining,salary,dept_name)
... VALUES(105,'Sarah','ASSOCIATE','2021-06-25',25000,'TESTING');
... APPLY BATCH;
cqlsh:employee> SELECT * FROM employee_info
... ;
emp_id | salary | date_of_joining | dept_name | designation | emp_name
-----+-----+-----+-----+-----+-----
105 | 25000 | 2021-06-24 18:30:00.000000+0000 | TESTING | ASSOCIATE | Sarah
111 | 25000 | 2021-06-21 18:30:00.000000+0000 | DEVELOPING | ASSOCIATE | Tom
121 | 35000 | 2021-03-29 18:30:00.000000+0000 | HR | MANAGER | Elsa
115 | 20000 | 2021-12-29 18:30:00.000000+0000 | DEVELOPING | ASSISTANT | Chris
100 | 30000 | 2021-09-10 18:30:00.000000+0000 | TESTING | MANAGER | John
(5 rows)
cqlsh:employee> UPDATE employee_info SET emp_name = 'Jessica', dept_name =
'DEVELOPING' WHERE
emp_id = 121;
cqlsh:employee> UPDATE employee_info SET emp_name = 'Jessica', dept_name =
'DEVELOPING' WHERE
emp_id = 121 AND salary = 35000;
cqlsh:employee> SELECT * FROM employee_info ;
emp_id | salary | date_of_joining | dept_name | designation | emp_name
-----+-----+-----+-----+-----+-----

```

```

105 | 25000 | 2021-06-24 18:30:00.000000+0000 | TESTING | ASSOCIATE | Sarah
111 | 25000 | 2021-06-21 18:30:00.000000+0000 | DEVELOPING | ASSOCIATE | Tom
121 | 35000 | 2021-03-29 18:30:00.000000+0000 | DEVELOPING | MANAGER | Jessica
115 | 20000 | 2021-12-29 18:30:00.000000+0000 | DEVELOPING | ASSISTANT | Chris
100 | 30000 | 2021-09-10 18:30:00.000000+0000 | TESTING | MANAGER | John

```

(5 rows)

```
cqlsh:employee> SELECT * FROM employee_info WHERE emp_id in (105, 111, 121, 115, 100) order by salary;
```

```
cqlsh:employee> paging off
```

Disabled Query paging.

```
cqlsh:employee> SELECT * FROM employee_info WHERE emp_id in (105, 111, 121, 115, 100) order by salary;
```

```
emp_id | salary | date_of_joining | dept_name | designation | emp_name
```

```

-----+-----+-----+-----+-----+-----
115 | 20000 | 2021-12-29 18:30:00.000000+0000 | DEVELOPING | ASSISTANT | Chris
105 | 25000 | 2021-06-24 18:30:00.000000+0000 | TESTING | ASSOCIATE | Sarah
111 | 25000 | 2021-06-21 18:30:00.000000+0000 | DEVELOPING | ASSOCIATE | Tom
100 | 30000 | 2021-09-10 18:30:00.000000+0000 | TESTING | MANAGER | John
121 | 35000 | 2021-03-29 18:30:00.000000+0000 | DEVELOPING | MANAGER | Jessica

```

(5 rows)

```
cqlsh:employee> ALTER TABLE employee_info ADD projects text;
```

```
cqlsh:employee> UPDATE employee_info SET projects = 'Chat App' WHERE emp_id = 111;
```

```
cqlsh:employee> UPDATE employee_info SET projects = 'Chat App' WHERE emp_id = 111 and salary = 25000;
```

```
cqlsh:employee> UPDATE employee_info SET projects = 'Discord Bot' WHERE emp_id = 115 and salary =
```

```
20000;
```

```
cqlsh:employee> UPDATE employee_info SET projects = 'Campus Portal' WHERE emp_id = 105 and salary =
```

```
25000;
```

```
cqlsh:employee> UPDATE employee_info SET projects = 'YouTube Downloader' WHERE  
emp_id = 100 and
```

```
salary = 30000;
```

```
cqlsh:employee> UPDATE employee_info SET projects = 'Library Management System '  
WHERE emp_id = 121
```

```
and salary = 35000;
```

```
cqlsh:employee> SELECT * FROM employee_infor
```

```
... ;
```

```
cqlsh:employee> SELECT * FROM employee_info ;
```

```
emp_id | salary | date_of_joining | dept_name | designation | emp_name | projects
```

```
-----+-----+-----+-----+-----+-----+-----  
-----
```

```
105 | 25000 | 2021-06-24 18:30:00.000000+0000 | TESTING | ASSOCIATE | Sarah | Campus  
Portal
```

```
111 | 25000 | 2021-06-21 18:30:00.000000+0000 | DEVELOPING | ASSOCIATE | Tom | Chat  
App
```

```
121 | 35000 | 2021-03-29 18:30:00.000000+0000 | DEVELOPING | MANAGER | Jessica |  
Library
```

```
Management System
```

```
115 | 20000 | 2021-12-29 18:30:00.000000+0000 | DEVELOPING | ASSISTANT | Chris |  
Discord
```

```
Bot
```

```
100 | 30000 | 2021-09-10 18:30:00.000000+0000 | TESTING | MANAGER | John | YouTube  
Downloader
```

```
(5 rows)
```

```
cqlsh:employee> INSERT INTO
```

```
... employee_info(emp_id,emp_name,designation,date_of_joining,salary,dept_name)
```

```
...
```

```
... ;
```

```
cqlsh:employee> INSERT INTO
```



```
... employee_info(emp_id,emp_name,designation,date_of_joining,salary,dept_name)
... VALUES(110,'SAM','ASSOCIATE','2021-01-11',28000,'TESTING') USING TTL 15;
cqlsh:employee> SELECT TTL(emp_name) from employee_info WHERE emp_id = 110;
ttl(emp_name)
```

```
-----
```

```
3
```

```
(1 rows)
```

```
cqlsh:employee> SELECT * FROM employee_info;
```

```
emp_id | salary | date_of_joining | dept_name | designation | emp_name | projects
```

```
-----+-----+-----+-----+-----+-----+-----
-----
```

```
105 | 25000 | 2021-06-24 18:30:00.000000+0000 | TESTING | ASSOCIATE | Sarah | Campus
Portal
```

```
111 | 25000 | 2021-06-21 18:30:00.000000+0000 | DEVELOPING | ASSOCIATE | Tom | Chat
App
```

```
121 | 35000 | 2021-03-29 18:30:00.000000+0000 | DEVELOPING | MANAGER | Jessica |
Library
```

```
Management System
```

```
115 | 20000 | 2021-12-29 18:30:00.000000+0000 | DEVELOPING | ASSISTANT | Chris |
Discord
```

```
Bot
```

```
100 | 30000 | 2021-09-10 18:30:00.000000+0000 | TESTING | MANAGER | John | YouTube
Downloader
```

```
(5 rows)
```

PROGRAM 2

Perform the following DB operations using Cassandra.

1.Create a keyspace by name Library

2. Create a column family by name Library-Info with attributes

Stud_Id Primary Key, Counter_value of type Counter,

Stud_Name, Book-Name, Book-Id, Date_of_issue

3. Insert the values into the table in batch

4. Display the details of the table created and increase the value of the counter

5. Write a query to show that a student with id 112 has taken a book “BDA” 2 times.

6. Export the created column to a csv file

7. Import a given csv dataset from local file system into Cassandra column family

```
cqlsh:library> CREATE KEYSPACE library WITH replication = {'class':  
'SimpleStrategy','replication_factor':1};
```

```
cqlsh:library> USE library ;
```

```
cqlsh:library> CREATE TABLE Library_info(stud_id int, stud_name text, book_name text,  
book_id text,
```

```
date_of_issue timestamp, counter_value counter, PRIMARY KEY(stud_id,stud_name,  
book_name, book_id,
```

```
date_of_issue));
```

```
cqlsh:library> BEGIN COUNTER BATCH
```

```
... UPDATE library_info set counter_value +=1 where stud_id = 111 and stud_name = 'Manoj'  
and
```

```
book_name = 'Operations Research' and book_id = '56TXT' and date_of_issue = '2021-09-12';
```

```
... UPDATE library_info set counter_value +=1 where stud_id = 112 and stud_name = 'Kamal'  
and
```

```
book_name = 'Engineering Mathematics-3' and book_id = '5ERW4' and date_of_issue =  
'2021-04-10';
```

```
... UPDATE library_info set counter_value +=1 where stud_id = 113 and stud_name = 'Mahesh'  
and
```

```
book_name = 'Robinson Crusoe' and book_id = '34EDC' and date_of_issue = '2021-02-01';
```

```
... UPDATE library_info set counter_value +=1 where stud_id = 114 and stud_name = 'Raj' and
```

```
book_name = 'Engineering Drawing' and book_id = '123ER' and date_of_issue = '2021-04-03';  
... APPLY BATCH;
```

```
cqlsh:library> SELECT * FROM library_info ;
```

```
stud_id | stud_name | book_name | book_id | date_of_issue | counter_value
```

```
-----+-----+-----+-----+-----+-----
```

```
114 | Raj | Engineering Drawing | 123ER | 2021-04-02 18:30:00.000000+0000 | 1
```

```
111 | Manoj | Operations Research | 56TXT | 2021-09-11 18:30:00.000000+0000 | 1
```

```
113 | Mahesh | Robinson Crusoe | 34EDC | 2021-01-31 18:30:00.000000+0000 | 1
```

```
112 | Kamal | Engineering Mathematics-3 | 5ERW4 | 2021-04-09 18:30:00.000000+0000 | 1
```

```
(4 rows)
```

```
cqlsh:library> UPDATE library_info set counter_value += 1 where stud_id = 112 and stud_name  
= 'Kamal' and
```

```
book_name = 'Engineering Mathematics-3' and book_id = '5ERW4' and date_of_issue =  
'2021-04-09';
```

```
cqlsh:library> SELECT * FROM library_info ;
```

```
stud_id | stud_name | book_name | book_id | date_of_issue | counter_value
```

```
-----+-----+-----+-----+-----+-----
```

```
114 | Raj | Engineering Drawing | 123ER | 2021-04-02 18:30:00.000000+0000 | 1
```

```
111 | Manoj | Operations Research | 56TXT | 2021-09-11 18:30:00.000000+0000 | 1
```

```
113 | Mahesh | Robinson Crusoe | 34EDC | 2021-01-31 18:30:00.000000+0000 | 1
```

```
112 | Kamal | Engineering Mathematics-3 | 5ERW4 | 2021-04-09 18:30:00.000000+0000 | 2
```

```
cqlsh:library> copy library_info(stud_id,stud_name, book_name, book_id,  
date_of_issue,counter_value) to
```

```
'library_info.csv' ;
```

```
Using 11 child processes
```

```
Starting copy of library.library_info with columns [stud_id, stud_name, book_name, book_id,  
date_of_issue,
```

```
counter_value].
```

```
Processed: 6 rows; Rate: 39 rows/s; Avg. rate: 39 rows/s
```

```
6 rows exported to 1 files in 0.165 seconds.
```

```
cqlsh:library> copy library_info(stud_id,stud_name, book_name, book_id,  
date_of_issue,counter_value) from
```

```
'library_info.csv' ;
```

Using 11 child processes

```
Starting copy of library.library_info with columns [stud_id, stud_name, book_name, book_id,  
date_of_issue,  
counter_value].
```

Processed: 6 rows; Rate: 10 rows/s; Avg. rate: 15 rows/s

6 rows imported from 1 files in 0.392 seconds (0 skipped).

PROGRAM 3

MongoDB CREATE DATABASE IN MONGODB.

```
> use myDB
switched to db myDB

db;
myDB

show dbs;
admin 0.000GB
config 0.000GB
local 0.000GB
```

II. CRUD (CREATE, READ, UPDATE, DELETE) OPERATIONS

1. To create a collection by the name “Student”. Let us take a look at the collection list prior to the creation of the new collection “Student”.

```
db.createCollection("Student"); => sql equivalent
CREATE TABLE STUDENT(...);
{ "ok" : 1 }
```

2.To drop a collection by the name “Student”.

```
db.Student.drop();
```

3.Create a collection by the name

“Students” and store the following data in it.

```
db.Student.insert({_id:1,StudName:"MichelleJacintha",Grade:"VII",Hobbies:"InternetSurfing"});
WriteResult({ "nInserted" : 1 })
```

4.Insert the document for “AryanDavid” in to the Students

collection only if it does not already exist in the collection.

However, if it is already present in the collection, then update the document with new values. (Update his Hobbies from “Skating” to

“Chess”.) Use “Update else insert” (if there is an existing document, it will attempt to update it, if there is no existing document then it will insert it).

```
db.Student.update({_id:3,StudName:"AryanDavid",Grade:" VII"},{$set:{Hobbies:"Skating"}},{upsert:true});
WriteResult({ "nMatched" : 0, "nUpserted" : 1, "nModified" : 0, "_id" : 3 })
```

5.FIND METHOD

A. To search for documents from the “Students” collection based on certain search criteria.

```
db.Student.find({StudName:"AryanDavid"});
({cond..},{columns.. column:1, columnname:0} )
{ "_id" : 3, "Grade" : "VII", "StudName" : "AryanDavid",
"Hobbies" : "Skating" }
```

B. To display only the StudName and Grade from all the documents of the Students collection. The identifier _id should be suppressed and NOT displayed.

```
db.Student.find({}, {StudName:1,Grade:1,_id:0});
{ "StudName" : "MichelleJacintha", "Grade" : "VII" }
{ "Grade" : "VII", "StudName" : "AryanDavid" }
```

C. To find those documents where the Grade is set to ‘VII’

```
db.Student.find({Grade:{$eq:'VII'}}).pretty();
{
  "_id" : 1,
  "StudName" : "MichelleJacintha",
  "Grade" : "VII",
  "Hobbies" : "InternetSurfing"
}
```

```
{
  "_id" : 3,
  "Grade" : "VII",
  "StudName" : "AryanDavid",
  "Hobbies" : "Skating"
}
```

D. To find those documents from the Students collection where the Hobbies is set to either ‘Chess’ or is set to ‘Skating’.

```
db.Student.find({Hobbies : { $in: ['Chess','Skating']}}).pretty()
{
  "_id" : 3,
  "Grade" : "VII",
  "StudName" : "AryanDavid",
  "Hobbies" : "Skating"
}
```

E. To find documents from the Students collection where the StudName begins with “M”.

```
db.Student.find({StudName:/^M/}).pretty();
{
  "_id" : 1,
  "StudName" : "MichelleJacintha",
  "Grade" : "VII",
  "Hobbies" : "InternetSurfing"
}
```

F. To find documents from the Students collection where the StudName has an “e” in any position.

```
db.Student.find({StudName:/e/}).pretty();
```

```
{
  "_id" : 1,
  "StudName" : "MichelleJacintha",
  "Grade" : "VII",
  "Hobbies" : "InternetSurfing"
}
```

G. To find the number of documents in the Students collection.

```
db.Student.count();
```

2

H. To sort the documents from the Students collection in the descending order of StudName.

```
db.Student.find().sort({StudName:-1}).pretty();
```

```
{
  "_id" : 1,
  "StudName" : "MichelleJacintha",
  "Grade" : "VII",
  "Hobbies" : "InternetSurfing"
}
{
  "_id" : 3,
  "Grade" : "VII",
  "StudName" : "AryanDavid",
  "Hobbies" : "Skating"
}
```

III. Import data from a CSV file

Given a CSV file “sample.txt” in the D:drive, import the file into the MongoDB collection, “SampleJSON”. The collection is in the database “test”.


```
mongoimport --db Student --collection airlines --type csv --
headerline --file /home/hduser/Desktop/airline.csv
```

IV. Export data to a CSV file

```
mongoexport --host localhost --db Student --collection
airlines --csv --out /home/hduser/Desktop/output.txt --
fields "Year"
```

```
,
```

```
"Quarter"
```

V. Save Method :

```
db.Student.save({StudName:"Vamsi", Grade:"VI"})
WriteResult({ "nInserted" : 1 })
```

VI. Add a new field to existing Document:

```
db.Student.update({_id:ObjectId("625695cc7d129fb98b44c8a1")},
{$set:{Location:"Network"}})
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
```

VII. Remove the field in an existing Document

```
db.Student.update({_id:ObjectId("625695cc7d129fb98b44c8a1
")},
{$unset:{Location:"Network"}})
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
```

VIII. Finding Document based on search criteria suppressing few fields

```
db.Student.find({_id:1},{StudName:1,Grade:1,_id:0});
{ "StudName" : "MichelleJacintha", "Grade" : "VII" }
```

To find those documents where the Grade is not set to 'VII'

```
db.Student.find({Grade:{$ne:'VII'}}).pretty();
{
  "_id" : ObjectId("625695cc7d129fb98b44c8a1"),
```

```
"StudName" : "Vamsi",  
"Grade" : "VI"  
}
```

To find documents from the Students collection where the StudName ends with s.

```
db.Student.find({StudName:/s$/}).pretty();  
  
{  
  "_id" : 1,  
  "StudName" : "MichelleJacintha",  
  "Grade" : "VII",  
  "Hobbies" : "InternetSurfing"  
}
```

IX. to set a particular field value to NULL

```
db.Student.update({_id:3},{ $set: {Location:null}})  
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
```

X. Count the number of documents in Student Collections

```
db.Student.count()  
  
3
```

**XI. Count the number of documents in Student Collections
with grade :VII**

```
db.Student.count({Grade:"VII"})
```

2 retrieve first 3 documents

```
db.Student.find({Grade:"VII"}).limit(1).pretty();  
  
{  
  "_id" : 1,  
  "StudName" : "MichelleJacintha",  
  "Grade" : "VII",  
  "Hobbies" : "InternetSurfing"
```

```
}
```

Sort the document in Ascending order

```
db.Student.find().sort({StudName:1}).pretty();
```

```
{
```

```
"_id" : 3,
```

```
"Grade" : "VII",
```

```
"StudName" : "AryanDavid",
```

```
"Hobbies" : "Skating",
```

```
"Location" : null
```

```
}
```

```
{
```

```
"_id" : 1,
```

```
"StudName" : "MichelleJacintha",
```

```
"Grade" : "VII",
```

```
"Hobbies" : "InternetSurfing"
```

```
}
```

```
{
```

```
"_id" : ObjectId("625695cc7d129fb98b44c8a1"),
```

```
"StudName" : "Vamsi",
```

```
"Grade" : "VI"
```

```
}
```

Note: for desending order :

```
db.Students.find().sort({StudName:-
```

```
1}).pretty();
```

To Skip the 1st two documents from the Students Collections

```
db.Student.find().skip(2).pretty()
```

```
{
```

```
"_id" : ObjectId("625695cc7d129fb98b44c8a1"),
```

```
"StudName" : "Vamsi",  
"Grade" : "VI"  
}
```

XII. Create a collection by name “food” and add to each document add a “fruits” array

```
db.food.insert( { _id:1, fruits:['grapes','mango','apple'] } )  
db.food.insert( { _id:2, fruits:['grapes','mango','cherry'] } )  
db.food.insert( { _id:3, fruits:['banana','mango'] } )  
{ "_id" : 1, "fruits" : [ "grapes", "mango", "apple" ] }  
{ "_id" : 2, "fruits" : [ "grapes", "mango", "cherry" ] }  
{ "_id" : 3, "fruits" : [ "banana", "mango" ] }
```

To find those documents from the “food” collection which has the “fruits array” consists of “grapes”, “mango” and “apple”.

```
db.food.find ( {fruits: ['grapes','mango','apple'] } ). pretty();  
{ "_id" : 1, "fruits" : [ "grapes", "mango", "apple" ] }
```

To find in “fruits” array having “mango” in the first index position.

```
db.food.find ( {“fruits.1”:grapes'} )
```

To find those documents from the “food” collection where the size of the array is two.

```
db.food.find ( {“fruits”: {$size:2}} )  
{ "_id" : 3, "fruits" : [ "banana", "mango" ] }
```

To find the document with a particular id and display the first two elements from the array “fruits”

```
db.food.find({_id:1},{“fruits”:$slice:2}})  
{ "_id" : 1, "fruits" : [ "grapes", "mango" ] }
```

To find all the documents from the food collection which have elements mango and grapes in the array “fruits”

```
db.food.find({fruits: {$all:[“mango”,”grapes”]}})  
{ "_id" : 1, "fruits" : [ "grapes", "mango", "apple" ] }  
{ "_id" : 2, "fruits" : [ "grapes", "mango", "cherry" ] }
```

Update on Array: using particular id replace the element present in the 1st index position of the fruits array with apple

```
db.food.update({_id:3},{ $set:{'fruits.1':'apple'}})
```

```
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
```

Insert new key value pairs in the fruits array

```
db.food.update({_id:2},{ $push:{price:{grapes:80,mango:200,cherry:100}}})
```

```
{ "_id" : 1, "fruits" : [ "grapes", "mango", "apple" ] }
```

```
{ "_id" : 2, "fruits" : [ "grapes", "mango", "cherry" ], "price" : [ {  
"grapes" : 80, "mango" : 200, "cherry" : 100 } ] }
```

```
{ "_id" : 3, "fruits" : [ "banana", "apple" ] }
```

PROGRAM 4

4. Screenshot of Hadoop installed

```
Apache Hadoop Distribution - yarn.nodemanager
INFO: Registering org.apache.hadoop.yarn.server.nodemanager.webapp.JAXBContextResolver as a provider class
Jul 15, 2022 9:47:49 PM com.sun.jersey.server.impl.application.WebApplicationImpl _initiate
INFO: Initiating Jersey application, version 'Jersey: 1.19.0/11/2015 03:25 AM'
Jul 15, 2022 9:47:49 PM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory getComponentProvider
INFO: Binding org.apache.hadoop.yarn.server.nodemanager.webapp.JAXBContextResolver to GuiceManagedComponentProvider with the scope 'Singleton'
Jul 15, 2022 9:47:49 PM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory getComponentProvider
INFO: Binding org.apache.hadoop.yarn.webapp.GenericExceptionHandler to GuiceManagedComponentProvider with the scope 'Singleton'
Jul 15, 2022 9:47:50 PM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory getComponentProvider
INFO: Binding org.apache.hadoop.yarn.server.nodemanager.webapp.NMWebServices to GuiceManagedComponentProvider with the scope 'Singleton'
2022-07-15 21:47:50,839 INFO Handler.ContextHandler: Started o.e.j.w.WebAppContext[/org.apache.hadoop.yarn.server.nodemanager.webapp-3.2.2-jar/webapps/node]
2022-07-15 21:47:50,848 INFO Server.Server: Started @8517ms
2022-07-15 21:47:50,849 INFO webapp.WebApps: Web app node started at 8042
2022-07-15 21:47:50,849 INFO nodemanager.NodeStatusUpdaterImpl: Node ID assigned is : DESKTOP-1GQKUN:50713
2022-07-15 21:47:50,859 INFO will.WaitPauseMonitor: Starting NM pause monitor
2022-07-15 21:47:50,853 INFO Client.RMPProxy: Connecting to ResourceManager at /0.0.0.0:8031
2022-07-15 21:47:50,180 INFO nodemanager.NodeStatusUpdaterImpl: Registering with RM using containers [1]
2022-07-15 21:47:50,391 INFO security.NMContainerTokenSecretManager: Rolling master-key for container-tokens, got key with id -37466286
2022-07-15 21:47:50,391 INFO security.WTTokenSecretManagerInRM: Rolling master-key for container-tokens, got key with id -695540823
2022-07-15 21:47:50,392 INFO nodemanager.NodeStatusUpdaterImpl: Registered with ResourceManager as DESKTOP-1GQKUN:50713 with total resource of <memory>:8192, <vcores>:8

Apache Hadoop Distribution - hadoop.datanode
2022-07-15 21:45:47,643 INFO Impl.FsDatasetImpl: Total time to scan all replicas for block pool BP-1951056128-192.168.1.7-1644143530070: 64ms
2022-07-15 21:45:47,644 INFO Impl.FsDatasetImpl: Adding replicas to map for block pool BP-1951056128-192.168.1.7-1644143530070 on volume Z:\BigData\hadoop-3.2.2\dfs\datanode...
2022-07-15 21:45:47,645 INFO Impl.BlockPoolSlice: Replica cache file: Z:\BigData\hadoop-3.2.2\dfs\datanode\current\BP-1951056128-192.168.1.7-1644143530070
2022-07-15 21:45:47,672 INFO Impl.FsDatasetImpl: Time to add replicas to map for block pool BP-1951056128-192.168.1.7-1644143530070 on volume Z:\BigData\hadoop-3.2.2\dfs\datanode: 27ms
2022-07-15 21:45:47,673 INFO Impl.FsDatasetImpl: Total time to add all replicas to map for block pool BP-1951056128-192.168.1.7-1644143530070: 29ms
2022-07-15 21:45:47,673 INFO checker.ThrottledAsyncChecker: Scheduling a check for Z:\BigData\hadoop-3.2.2\dfs\datanode
2022-07-15 21:45:47,682 INFO checker.DatasetVolumeChecker: Scheduled health check for volume Z:\BigData\hadoop-3.2.2\dfs\datanode
2022-07-15 21:45:47,701 INFO datanode.VolumeScanner: VolumeScanner(Z:\BigData\hadoop-3.2.2\dfs\datanode, DS-6d35dabc-0008-4c85-8303-e3005e89107c) no suitable block pools found to scan. Waiting 608433460 ms.
2022-07-15 21:45:47,733 INFO datanode.DirectoryScanner: Periodic Directory Tree Verification scan starting at 16/07/22, 1:31 AM with interval of 2160000ms
2022-07-15 21:45:47,741 INFO datanode.Datanode: Block pool BP-1951056128-192.168.1.7-1644143530070 (Datanode Uuid 7c7da9eb-c504-4f5f-abbe-de408c4286cc) service to localhost/127.0.0.1:9866 beginning handshake with NM
2022-07-15 21:45:47,767 INFO datanode.Datanode: Block pool BP-1951056128-192.168.1.7-1644143530070 (Datanode Uuid 7c7da9eb-c504-4f5f-abbe-de408c4286cc) service to localhost/127.0.0.1:9866 successfully registered with NM
2022-07-15 21:45:47,787 INFO datanode.Datanode: For namode: localhost/127.0.0.1:9866 using BLOCKREPORT_INTERVAL of 2160000ms (CACHEREPORT_INTERVAL of 10000ms Initial delay: msec; heartbeatInterval:3000)
2022-07-15 21:45:47,945 INFO datanode.Datanode: Successfully sent block report 0x6ce2a69a720089c0, containing 1 storage report(s), of which we sent 1. The reports had 0 total blocks and used 1 RPC(s). This took 3 msec to generate and 50 msec for RPC and NM processing. Got back one command: FinalizeCommand/5
2022-07-15 21:45:47,946 INFO datanode.Datanode: Got finalize command for block pool BP-1951056128-192.168.1.7-1644143530070

Apache Hadoop Distribution - hadoop.nameservice
at org.apache.hadoop.hdfs.protocol.proto.DatanodeProtocolProtos$DatanodeProtocolService$2.callBlockingMethod(DatanodeProtocolProtos.java:31672)
at org.apache.hadoop.ipc.ProtobufRpcEngine$Server$ProtoBufRpcInvoker.call(ProtobufRpcEngine.java:528)
at org.apache.hadoop.ipc.RPC$Server.call(RPC.java:1086)
at org.apache.hadoop.ipc.Server$RPC$Call.run(Server.java:1029)
at org.apache.hadoop.ipc.Server$RPC$Call.run(Server.java:957)
at java.base/java.security.AccessController.doPrivileged(Native Method)
at java.base/java.security.auth.Subject.doAs(Subject.java:423)
at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1762)
at org.apache.hadoop.ipc.Server$Handler.run(Server.java:2957)
2022-07-15 21:45:48,588 INFO fs.TrashPolicyDefault: The configured checkpoint interval is 0 minutes. Using an interval of 1440 minutes that is used for deletion instead
2022-07-15 21:45:48,588 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 1440 minutes, Eruption interval = 1440 minutes.
2022-07-15 21:45:47,775 INFO hdfs.StateChanger: BLOCK* registerDatanode: from DatanodeRegistration(127.0.0.1:9866, datanodeuid=7c7da9eb-c504-4f5f-abbe-de408c4286cc, infoPort=9866, infoSecurePort=0, ipcPort=9867, storageInfo=lv=57;cid=CID-0761f835-5479-48e3-a70d-1b47acc13d6f;nsid=557392828;c=1644143530070) storage 7c7da9eb-c504-4f5f-abbe-de408c4286cc
2022-07-15 21:45:47,868 INFO net.NetworkTopology: Adding a new node: /default-rack/127.0.0.1:9866
2022-07-15 21:45:47,776 INFO Blockmanagement.BlockReportLeaseManager: Registered DN 7c7da9eb-c504-4f5f-abbe-de408c4286cc (127.0.0.1:9866)
2022-07-15 21:45:47,868 INFO Blockmanagement.DatanodeDescriptor: Adding new storage ID DS-6d35dabc-0008-4c85-8303-e3005e89107c for DN 127.0.0.1:9866
2022-07-15 21:45:47,900 INFO BlockStateChanger: BLOCK* processReport 0x6ce2a69a720089c0: Processing first storage report for DS-6d35dabc-0008-4c85-8303-e3005e89107c from datanode 7c7da9eb-c504-4f5f-abbe-de408c4286cc
2022-07-15 21:45:47,900 INFO BlockStateChanger: BLOCK* processReport 0x6ce2a69a720089c0: From storage DS-6d35dabc-0008-4c85-8303-e3005e89107c node DatanodeRegistration(127.0.0.1:9866, datanodeuid=7c7da9eb-c504-4f5f-abbe-de408c4286cc, infoPort=9866, infoSecurePort=0, ipcPort=9867, storageInfo=lv=57;cid=CID-0761f835-5479-48e3-a70d-1b47acc13d6f;nsid=557392828;c=1644143530070), blocks: 0, hasStaleStorage: false, processing time: 2 msec, invalidatedBlocks: 0
```

```
Administrator: Command Prompt

Microsoft Windows [Version 10.0.22000.795]
(c) Microsoft Corporation. All rights reserved.

C:\WINDOWS\system32>start-dfs
'C:\Program' is not recognized as an internal or external command,
operable program or batch file.

C:\WINDOWS\system32>start-yarn
starting yarn daemons
'C:\Program' is not recognized as an internal or external command,
operable program or batch file.

C:\WINDOWS\system32>jps
12080 NameNode
14708 ResourceManager
2232 Jps
2648 DataNode
14556 NameManager

C:\WINDOWS\system32>
```

PROGRAM 5

Execution of HDFS Commands for interaction with Hadoop

Environment. (Minimum 10 commands to be executed)

start-all.sh

This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh

Starting namenodes on [localhost]

hduser@localhost's password:

localhost: starting namenode, logging to

/usr/local/hadoop/logs/hadoop-hduser-namenode-bmsce-Precision-T1700.out

hduser@localhost's password:

localhost: starting datanode, logging to

/usr/local/hadoop/logs/hadoop-hduser-datanode-bmsce-Precision-T1700.out

Starting secondary namenodes [0.0.0.0]

hduser@0.0.0.0's password:

0.0.0.0: starting secondarynamenode, logging to

/usr/local/hadoop/logs/hadoop-hduser-secondarynamenode-bmsce-Precision-T1700.out

starting yarn daemons

starting resourcemanager, logging to

/usr/local/hadoop/logs/yarn-hduser-resourcemanager-bmsce-Precision-T1700.out

hduser@localhost's password:

localhost: starting nodemanager, logging to

/usr/local/hadoop/logs/yarn-hduser-nodemanager-bmsce-Precision-T1700.out

hduser@bmsce-Precision-T1700:~\$ jps

5072 SecondaryNameNode

4674 NameNode

4856 DataNode

5563 NodeManager

6507 Jps

5231 ResourceManager

hduser@bmsce-Precision-T1700:~\$ ls

derby.log	Pictures
Desktop	pig_1564816082257.log
Documents	pig_1599215374374.log
Downloads	pt
examples.desktop	PT72Installer
first.text	Public
hadoop-2.6.0.tar.gz	R
hive	TCPclient.py
lol	TCPserver.py
metastore_db	Templates
Music	toinstalledlist
newfile.txt	UDPclient.py
newnewfile.txt	UDPserver.py

'Packet Tracer 7.2.1 for Linux 64 bit.tar.gz' Videos

```
hduser@bmsce-Precision-T1700:~$ /home/desktop
bash: /home/desktop: No such file or directory
hduser@bmsce-Precision-T1700:~$ /home/Desktop
bash: /home/Desktop: No such file or directory
hduser@bmsce-Precision-T1700:~$ cd desktop
bash: cd: desktop: No such file or directory
hduser@bmsce-Precision-T1700:~$ cd Desktop
hduser@bmsce-Precision-T1700:~/Desktop$ cd ..
hduser@bmsce-Precision-T1700:~$ hdfs dfs -mkdir ~/Desktop/Ishan
mkdir: `/home/hduser/Desktop/Ishan': No such file or directory
hduser@bmsce-Precision-T1700:~$ hdfs dfs -mkdir /abc
hduser@bmsce-Precision-T1700:~$ ls
```

derby.log	Pictures
Desktop	pig_1564816082257.log

Documents pig_1599215374374.log

Downloads pt

examples.desktop PT72Installer

first.text Public

hadoop-2.6.0.tar.gz R

hive TCPclient.py

lol TCPserver.py

metastore_db Templates

Music toinstalledlist

newfile.txt UDPclient.py

newnewfile.txt UDPserver.py

'Packet Tracer 7.2.1 for Linux 64 bit.tar.gz' Videos

hduser@bmsce-Precision-T1700:~\$ hdfs dfs -ls /

Found 7 items

drwxr-xr-x - hduser supergroup 0 2022-06-03 12:52 /SharmaJi

drwxr-xr-x - hduser supergroup 0 2022-06-04 09:34 /abc

drwxr-xr-x - hduser supergroup 0 2022-06-03 15:44 /bhavana

drwxr-xr-x - hduser supergroup 0 2022-06-01 15:22 /lochan

drwxr-xr-x - hduser supergroup 0 2022-06-03 15:45 /u1

-rw-r--r-- 1 hduser supergroup 19 2022-05-31 11:01 /user

drwxr-xr-x - hduser supergroup 0 2022-06-01 10:08 /vallisha

hduser@bmsce-Precision-T1700:~\$ cat newfile.txt

SharmaJi

KhanwaJi

PaiJI

Kasturba

pandeyji

patilwa

Nairwa

```
hduser@bmsce-Precision-T1700:~$ hdfs dfs -put /home/hduser/newfile.txt /abc/ishan.txt
```

```
hduser@bmsce-Precision-T1700:~$ hdfs dfs cat /abc/ishan.txt
```

cat: Unknown command

Did you mean -cat? This command begins with a dash.

```
hduser@bmsce-Precision-T1700:~$ hdfs dfs -cat /abc/ishan.txt
```

SharmaJi

KhanwaJi

PaiJI

Kasturba

pandeyji

patilwa

Nairwa

```
hduser@bmsce-Precision-T1700:~$ cat newnewfile.txt
```

SharmaJi

KhanwaJi

PaiJI

Kasturba

pandeyji

patilwa

Nairwa

```
hduser@bmsce-Precision-T1700:~$ cat > ishan.txt
```

Hello

This is a new text file

^C

```
hduser@bmsce-Precision-T1700:~$ cat ishan.txt
```

Hello

This is a new text file

```
hduser@bmsce-Precision-T1700:~$ hdfs dfs -copyFromLocal /home/hduser/ishan.txt /abc/ishan2.txt
```

```
hduser@bmsce-Precision-T1700:~$ hdfs dfs -cat /abc/ishan2.txt
```

Hello

This is a new text file

```
hduser@bmsce-Precision-T1700:~$ hdfs dfs -get /abc/ishan2.txt /home/hduser/ishan2_copy.txt
```

```
hduser@bmsce-Precision-T1700:~$ ls
```

derby.log 'Packet Tracer 7.2.1 for Linux 64 bit.tar.gz'

Desktop Pictures

Documents pig_1564816082257.log

Downloads pig_1599215374374.log

examples.desktop pt

first.text PT72Installer

hadoop-2.6.0.tar.gz Public

hive R

ishan2_copy.txt TCPclient.py

ishan.txt TCPserver.py

lol Templates

metastore_db toinstalledlist

Music UDPclient.py

newfile.txt UDPserver.py

newnewfile.txt Videos

```
hduser@bmsce-Precision-T1700:~$ hdfs dfs -getmerge /abc/ishan.txt /abc/ishan2.txt /home/hduser/ishan_merge.txt
```

```
hduser@bmsce-Precision-T1700:~$ ls
```

derby.log 'Packet Tracer 7.2.1 for Linux 64 bit.tar.gz'

Desktop Pictures

Documents pig_1564816082257.log

Downloads pig_1599215374374.log

examples.desktop pt
first.text PT72Installer
hadoop-2.6.0.tar.gz Public
hive R
ishan2_copy.txt TCPclient.py
ishan_merge.txt TCPserver.py
ishan.txt Templates
lol toinstalledlist
metastore_db UDPclient.py
Music UDPserver.py
newfile.txt Videos
newnewfile.txt

hduser@bmsce-Precision-T1700:~\$ ishan_merge.txt

ishan_merge.txt: command not found

hduser@bmsce-Precision-T1700:~\$ cat ishan_merge.txt

SharmaJi

KhanwaJi

PaiJI

Kasturba

pandeyji

patilwa

Nairwa

Hello

This is a new text file

hduser@bmsce-Precision-T1700:~\$ hadoop fs -getfacl /abc/

file: /abc

owner: hduser

group: supergroup

user::rwx

group::r-x

other::r-x

hduser@bmsce-Precision-T1700:~\$ hdfs dfs -copyToLocal /abc/ishan.txt /home/hduser/Desktop

hduser@bmsce-Precision-T1700:~\$ hdfs dfs -ls /

Found 7 items

drwxr-xr-x - hduser supergroup 0 2022-06-03 12:52 /SharmaJi

drwxr-xr-x - hduser supergroup 0 2022-06-04 09:40 /abc

drwxr-xr-x - hduser supergroup 0 2022-06-03 15:44 /bhavana

drwxr-xr-x - hduser supergroup 0 2022-06-01 15:22 /lochan

drwxr-xr-x - hduser supergroup 0 2022-06-03 15:45 /u1

-rw-r--r-- 1 hduser supergroup 19 2022-05-31 11:01 /user

drwxr-xr-x - hduser supergroup 0 2022-06-01 10:08 /vallisha

hduser@bmsce-Precision-T1700:~\$ hdfs dfs -mkdir /ishan

hduser@bmsce-Precision-T1700:~\$ hadoop fs -mv /abc /ishan

hduser@bmsce-Precision-T1700:~\$ hdfs dfs -ls /ishan

Found 1 items

drwxr-xr-x - hduser supergroup 0 2022-06-04 09:40 /ishan/abc

hduser@bmsce-Precision-T1700:~\$ hdfs dfs -ls /abc

ls: `/abc': No such file or directory

hduser@bmsce-Precision-T1700:~\$ hdfs dfs -ls /

Found 7 items

drwxr-xr-x - hduser supergroup 0 2022-06-03 12:52 /SharmaJi

drwxr-xr-x - hduser supergroup 0 2022-06-03 15:44 /bhavana

drwxr-xr-x - hduser supergroup 0 2022-06-04 09:59 /ishan

drwxr-xr-x - hduser supergroup 0 2022-06-01 15:22 /lochan

drwxr-xr-x - hduser supergroup 0 2022-06-03 15:45 /u1

```

-rw-r--r-- 1 hduser supergroup      19 2022-05-31 11:01 /user
drwxr-xr-x - hduser supergroup      0 2022-06-01 10:08 /vallisha
hduser@bmsce-Precision-T1700:~$ hdfs dfs -cat/ishan.txt
-cat/ishan.txt: Unknown command
hduser@bmsce-Precision-T1700:~$ hdfs dfs -cat /ishan.txt
cat: `/ishan.txt': No such file or directory
hduser@bmsce-Precision-T1700:~$ hdfs dfs -cat /ishan/ishan.txt
cat: `/ishan/ishan.txt': No such file or directory
hduser@bmsce-Precision-T1700:~$ hdfs dfs -ls /vallisha
Found 1 items
-rw-r--r-- 1 hduser supergroup      13 2022-06-01 09:52 /vallisha/sample1.txt
hduser@bmsce-Precision-T1700:~$ hdfs dfs -ls /ishan/abc
Found 2 items
-rw-r--r-- 1 hduser supergroup      57 2022-06-04 09:37 /ishan/abc/ishan.txt
-rw-r--r-- 1 hduser supergroup      30 2022-06-04 09:40 /ishan/abc/ishan2.txt
hduser@bmsce-Precision-T1700:~$ hdfs dfs -cat /ishan/ishan.txt
cat: `/ishan/ishan.txt': No such file or directory
hduser@bmsce-Precision-T1700:~$ hdfs dfs -cat /ishan/abc/ishan.txt
SharmaJi
KhanwaJi
PaiJI
Kasturba
pandeyji
patilwa
Nairwa
hduser@bmsce-Precision-T1700:~$ hdfs dfs -cp /vallisha/sample1.txt /ishan
hduser@bmsce-Precision-T1700:~$ hdfs dfs -ls /ishan
Found 2 items

```

```

drwxr-xr-x - hduser supergroup      0 2022-06-04 09:40 /ishan/abc
-rw-r--r-- 1 hduser supergroup     13 2022-06-04 10:07 /ishan/sample1.txt
hduser@bmsce-Precision-T1700:~$ hdfs dfs -cat /ishan/sample1.txt
sample1 text
hduser@bmsce-Precision-T1700:~$ sudo nano new_file.txt
[sudo] password for hduser:
hduser@bmsce-Precision-T1700:~$ ls
derby.log          newnewfile.txt
Desktop            'Packet Tracer 7.2.1 for Linux 64 bit.tar.gz'
Documents          Pictures
Downloads          pig_1564816082257.log
examples.desktop   pig_1599215374374.log
first.text         pt
hadoop-2.6.0.tar.gz PT72Installer
hive               Public
ishan2_copy.txt    R
ishan_merge.txt    TCPclient.py
ishan.txt          TCPserver.py
lol               Templates
metastore_db       toinstalledlist
Music              UDPclient.py
new_file.txt       UDPserver.py
newfile.txt        Videos
hduser@bmsce-Precision-T1700:~$ cat new_file.txt
This is a new file, created using sudo nano

```

PROGRAM 6

From the following link extract the weather data

<https://github.com/tomwhite/hadoop-book/tree/master/input/ncdc/all>

Create a Map Reduce program to

a) find average temperature for each year from NCDC data set.

\\AVERAGE DRIVER

```
package temp;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class AverageDriver {

    public static void main(String[] args) throws Exception {
        if (args.length != 2) {

            System.err.println("Please Enter the input and output parameters");
            System.exit(-1);
        }

        Job job = new Job();
        job.setJarByClass(AverageDriver.class);
        job.setJobName("Max temperature");
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        job.setMapperClass(AverageMapper.class);
        job.setReducerClass(AverageReducer.class);
        job.setOutputKeyClass(Text.class);
```



```

job.setOutputValueClass(IntWritable.class);
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

\\AVERAGE MAPPER

```

package temp;

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class AverageMapper extends Mapper<LongWritable, Text, Text, IntWritable> {
    public static final int MISSING = 9999;

    public void map(LongWritable key, Text value, Mapper<LongWritable, Text, Text,
        IntWritable>.Context context) throws IOException, InterruptedException {
        int temperature;

        String line = value.toString();
        String year = line.substring(15, 19);
        if (line.charAt(87) == '+') {
            temperature = Integer.parseInt(line.substring(88, 92));
        } else {
            temperature = Integer.parseInt(line.substring(87, 92));
        }

        String quality = line.substring(92, 93);
        if (temperature != 9999 && quality.matches("[01459]"))
            context.write(new Text(year), new IntWritable(temperature));
    }
}

```

```
}
```

\\AVERAGE REDUCER

```
package temp;

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class AverageReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values, Reducer<Text, IntWritable,
    Text, IntWritable>.Context context) throws IOException, InterruptedException {
        int max_temp = 0;
        int count = 0;
        for (IntWritable value : values) {
            max_temp += value.get();
            count++;
        }
        context.write(key, new IntWritable(max_temp / count));
    }
}
```



A terminal window showing the command `c:\hadoop_new\sbin>hdfs dfs -cat /tempAverageOutput/part-r-00000` and its output, which is a table of years and average temperatures.

1901	46
1949	94
1950	3

b) find the mean max temperature for every month

\\MEAN MAX DRIVER

```
package meanmax;
```

```

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;

import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
public class MeanMaxDriver {
public static void main(String[] args) throws Exception {
if (args.length != 2) {
System.err.println("Please Enter the input and output parameters");
System.exit(-1);
}
Job job = new Job();
job.setJarByClass(MeanMaxDriver.class);
job.setJobName("Max temperature");
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
job.setMapperClass(MeanMaxMapper.class);
job.setReducerClass(MeanMaxReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

\\MEAN MAX MAPPER
package meanmax;

```

```

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class MeanMaxMapper extends Mapper<LongWritable, Text, Text, IntWritable> {
    public static final int MISSING = 9999;

    public void map(LongWritable key, Text value, Mapper<LongWritable, Text, Text,
        IntWritable>.Context context) throws IOException, InterruptedException {
        int temperature;

        String line = value.toString();
        String month = line.substring(19, 21);
        if (line.charAt(87) == '+') {
            temperature = Integer.parseInt(line.substring(88, 92));
        } else {
            temperature = Integer.parseInt(line.substring(87, 92));
        }

        String quality = line.substring(92, 93);
        if (temperature != 9999 && quality.matches("[01459]"))
            context.write(new Text(month), new IntWritable(temperature));
    }
}

```

\\MEAN MAX REDUCER

```

package meanmax;

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;

```

```

import org.apache.hadoop.mapreduce.Reducer;

public class MeanMaxReducer extends Reducer<Text, IntWritable, Text, IntWritable> {

    public void reduce(Text key, Iterable<IntWritable> values, Reducer<Text, IntWritable,
    Text, IntWritable>.Context context) throws IOException, InterruptedException {

        int max_temp = 0;
        int total_temp = 0;
        int count = 0;

        int days = 0;
        for (IntWritable value : values) {
            int temp = value.get();
            if (temp > max_temp)
                max_temp = temp;
            count++;
            if (count == 3) {
                total_temp += max_temp;
                max_temp = 0;
                count = 0;
                days++;
            }
        }
        context.write(key, new IntWritable(total_temp / days));
    }
}

```

```
c:\hadoop_new\sbin>hdfs dfs -cat /tempMaxOutput/part-r-00000
01      44
02      17
03     111
04     194
05     256
06     278
07     317
08     283
09     211
10     156
11      89
12     117
```

PROGRAM 7

For a given Text file, Create a Map Reduce program to sort the content in an alphabetic order listing only top 10 maximum occurrences of words.

\\TOPN

```
package samples.topn;

import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.GenericOptionsParser;

public class TopN {

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        String[] otherArgs = (new GenericOptionsParser(conf, args)).getRemainingArgs();
        if (otherArgs.length != 2) {
            System.err.println("Usage: TopN <in> <out>");
            System.exit(2);
        }
        Job job = Job.getInstance(conf);
        job.setJobName("Top N");
        job.setJarByClass(TopN.class);
        job.setMapperClass(TopNMapper.class);
```

```

job.setReducerClass(TopNReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
FileInputFormat.addInputPath(job, new Path(otherArgs[0]));
FileOutputFormat.setOutputPath(job, new Path(otherArgs[1]));
System.exit(job.waitForCompletion(true) ? 0 : 1);
}

```

\\TOPN COMBINER

```

package samples.topn;

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class TopNCombiner extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values, Reducer<Text, IntWritable,
    Text, IntWritable>.Context context) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values)
            sum += val.get();
        context.write(key, new IntWritable(sum));
    }
}

```

\\TOPN MAPPER

```

package samples.topn;

import java.io.IOException;
import java.util.StringTokenizer;

```



```

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class TopNMapper extends Mapper<Object, Text, Text, IntWritable> {
    private static final IntWritable one = new IntWritable(1);
    private Text word = new Text();
    private String tokens = "[_#$%<>\\^=\\[\\]\\|\\*\\/\\\\\\,;\\.\\|-:()?!\"'"]";

    public void map(Object key, Text value, Mapper<Object, Text, Text, IntWritable>.Context
    context) throws IOException, InterruptedException {
        String cleanLine = value.toString().toLowerCase().replaceAll(this.tokens, " ");
        StringTokenizer itr = new StringTokenizer(cleanLine);
        while (itr.hasMoreTokens()) {
            this.word.set(itr.nextToken().trim());
            context.write(this.word, one);
        }
    }
}

```

\\TOPN REDUCER

```

package samples.topn;

import java.io.IOException;
import java.util.HashMap;
import java.util.Map;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;
import utils.MiscUtils;

public class TopNReducer extends Reducer<Text, IntWritable, Text, IntWritable> {

```

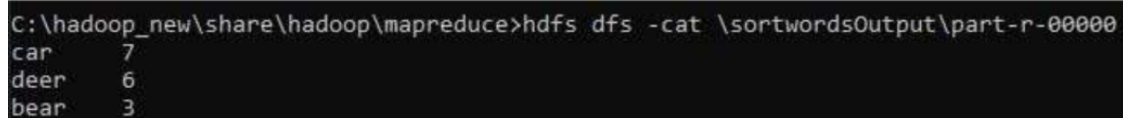
```

private Map<Text, IntWritable> countMap = new HashMap<>();

public void reduce(Text key, Iterable<IntWritable> values, Reducer<Text, IntWritable,
Text, IntWritable>.Context context) throws IOException, InterruptedException {
    int sum = 0;
    for (IntWritable val : values)
        sum += val.get();
    this.countMap.put(new Text(key), new IntWritable(sum));
}

protected void cleanup(Reducer<Text, IntWritable, Text, IntWritable>.Context context)
throws IOException, InterruptedException {
    Map<Text, IntWritable> sortedMap = MiscUtils.sortByValues(this.countMap);
    int counter = 0;
    for (Text key : sortedMap.keySet()) {
        if (counter++ == 20)
            break;
        context.write(key, sortedMap.get(key));
    }
}
}

```



```

C:\hadoop_new\share\hadoop\mapreduce>hdfs dfs -cat \sortwordsOutput\part-r-00000
car      7
deer     6
bear     3

```

PROGRAM 8

Create a Map Reduce program to demonstrating join operation

```
// JoinDriver.java

import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;
import org.apache.hadoop.mapred.libMultipleInputs;
import org.apache.hadoop.util.*;

public class JoinDriver extends Configured implements Tool {
    public static class KeyPartitioner implements Partitioner<TextPair, Text> {
        @Override

        public void configure(JobConf job) {}

        @Override
        public int getPartition(TextPair key, Text value, int numPartitions) {
            return (key.getFirst().hashCode() & Integer.MAX_VALUE) %
                numPartitions;
        }
    }

    @Override
    public int run(String[] args) throws Exception {
        if (args.length != 3) {
            System.out.println("Usage: <Department Emp Strength input>
            <Department Name input> <output>");
            return -1;
        }
        JobConf conf = new JobConf(getConf(), getClass());
```

```

conf.setJobName("Join 'Department Emp Strength input' with 'Department Name
input");
Path AInputPath = new Path(args[0]);
Path BInputPath = new Path(args[1]);
Path outputPath = new Path(args[2]);
MultipleInputs.addInputPath(conf, AInputPath, TextInputFormat.class,
Posts.class);
MultipleInputs.addInputPath(conf, BInputPath, TextInputFormat.class,
User.class);
FileOutputFormat.setOutputPath(conf, outputPath);
conf.setPartitionerClass(KeyPartitioner.class);
conf.setOutputValueGroupingComparator(TextPair.FirstComparator.class);
conf.setMapOutputKeyClass(TextPair.class);
conf.setReducerClass(JoinReducer.class);
conf.setOutputKeyClass(Text.class);

JobClient.runJob(conf);
return 0;
}

public static void main(String[] args) throws Exception {
int exitCode = ToolRunner.run(new JoinDriver(), args);
System.exit(exitCode);
}
}

// JoinReducer.java
import java.io.IOException;
import java.util.Iterator;

```

```

import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;

public class JoinReducer extends MapReduceBase implements Reducer<TextPair, Text,
Text,
Text> {
    @Override
    public void reduce (TextPair key, Iterator<Text> values, OutputCollector<Text, Text>
output, Reporter reporter)
        throws IOException
    {
        Text nodeId = new Text(values.next());
        while (values.hasNext()) {
            Text node = values.next();
            Text outValue = new Text(nodeId.toString() + "\t\t" + node.toString());
            output.collect(key.getFirst(), outValue);
        }
    }
}

```

//Posts.java

```

import java.io.IOException;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;

public class Posts extends MapReduceBase implements Mapper<LongWritable, Text,
TextPair,
Text> {
    @Override
    public void map(LongWritable key, Text value, OutputCollector<TextPair, Text> output,

```

```

Reporter reporter)
throws IOException
{
String valueString = value.toString();
String[] SingleNodeData = valueString.split("\t");
output.collect(new TextPair(SingleNodeData[3], "0"), new
Text(SingleNodeData[9]));
}
}

```

```

// TextPair.java
import java.io.*;
import org.apache.hadoop.io.*;

public class TextPair implements WritableComparable<TextPair> {
private Text first;
private Text second;
public TextPair() {
set(new Text(), new Text());
}
public TextPair(String first, String second) {
set(new Text(first), new Text(second));
}
public TextPair(Text first, Text second) {
set(first, second);
}
public void set(Text first, Text second) {
this.first = first;
this.second = second;
}
}

```

```

    }

    public Text getFirst() {
        return first;
    }

    public Text getSecond() {
        return second;
    }

    @Override
    public void write(DataOutput out) throws IOException {
        first.write(out);
        second.write(out);
    }

    @Override
    public void readFields(DataInput in) throws IOException {
        first.readFields(in);
        second.readFields(in);
    }

    @Override
    public int hashCode() {

        return first.hashCode() * 163 + second.hashCode();
    }

    @Override
    public boolean equals(Object o) {
        if (o instanceof TextPair) {
            TextPair tp = (TextPair) o;
            return first.equals(tp.first) && second.equals(tp.second);
        }
    }

```

```

return false;
}

@Override
public String toString() {
return first + "\t" + second;
}

@Override
public int compareTo(TextPair tp) {
int cmp = first.compareTo(tp.first);
if (cmp != 0) {
return cmp;
}
return second.compareTo(tp.second);
}

// ^^ TextPair
// vv TextPairComparator
public static class Comparator extends WritableComparator {
private static final Text.Comparator TEXT_COMPARATOR = new Text.Comparator();
public Comparator() {
super(TextPair.class);
}

@Override
public int compare(byte[] b1, int s1, int l1,
byte[] b2, int s2, int l2) {
try {
int firstL1 = WritableUtils.decodeVIntSize(b1[s1]) + readVInt(b1, s1);
int firstL2 = WritableUtils.decodeVIntSize(b2[s2]) + readVInt(b2, s2);
int cmp = TEXT_COMPARATOR.compare(b1, s1, firstL1, b2, s2, firstL2);

```



```

if (cmp != 0) {
return cmp;
}
return TEXT_COMPARATOR.compare(b1, s1 + firstL1, l1 - firstL1,

b2, s2 + firstL2, l2 - firstL2);
} catch (IOException e) {
throw new IllegalArgumentException(e);
}
}
}
static {
WritableComparator.define(TextPair.class, new Comparator());
}
public static class FirstComparator extends WritableComparator {
private static final Text.Comparator TEXT_COMPARATOR = new Text.Comparator();
public FirstComparator() {
super(TextPair.class);
}
@Override
public int compare(byte[] b1, int s1, int l1,
byte[] b2, int s2, int l2) {
try {
int firstL1 = WritableUtils.decodeVIntSize(b1[s1]) + readVInt(b1, s1);
int firstL2 = WritableUtils.decodeVIntSize(b2[s2]) + readVInt(b2, s2);
return TEXT_COMPARATOR.compare(b1, s1, firstL1, b2, s2, firstL2);
} catch (IOException e) {
throw new IllegalArgumentException(e);
}
}
}

```

```

    }
    }

    @Override
    public int compare(WritableComparable a, WritableComparable b) {
        if (a instanceof TextPair && b instanceof TextPair) {
            return ((TextPair) a).first.compareTo(((TextPair) b).first);
        }
        return super.compare(a, b);
    }
} }

// User.java
import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FSDataInputStream;
import org.apache.hadoop.fs.FSDataOutputStream;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;

import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;
import org.apache.hadoop.io.IntWritable;

public class User extends MapReduceBase implements Mapper<LongWritable, Text,
TextPair,
Text> {
    @Override

```

```

public void map(LongWritable key, Text value, OutputCollector<TextPair, Text> output,
Reporter reporter)
throws IOException
{
String valueString = value.toString();
String[] SingleNodeData = valueString.split("\t");
output.collect(new TextPair(SingleNodeData[0], "1"), new
Text(SingleNodeData[1]));
}
}

```

```

hduser@bmsce-Precision-T1700:~$ hadoop fs -cat /output_mapreduce/*
A11      50      Finance
B12      100     HR
C13      250     Manufacturing
Dept_ID  Total_Employee  Dept_Name
hduser@bmsce-Precision-T1700:~$

```

PROGRAM 9

Program to print word count on scala shell and print “Hello world” on scala IDE

```
val data=sc.textFile("sparkdata.txt")
data.collect;
val splitdata = data.flatMap(line => line.split(" "));
splitdata.collect;
val mapdata = splitdata.map(word => (word,1));
mapdata.collect;
val reducedata = mapdata.reduceByKey(_+_);
reducedata.collect;
```

```
package hellooWorld
```

```
object hello {
```

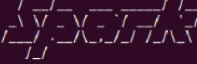
```
  def main (args: Array[String]) {
```

```
    println("Hello World")
```

```
  }
```

```
}
```

```
Spark session available as 'spark'.
Welcome to

 version 2.4.8

Using Scala version 2.11.12 (OpenJDK 64-Bit Server VM, Java 1.8.0_232)
Type in expressions to have them evaluated.
Type :help for more information.

scala> val data=sc.textFile("sample.txt")
data: org.apache.spark.rdd.RDD[String] = Sample.txt MapPartitionsRDD[1] at textFile at <console>:24

scala> data.collect;
res0: Array[String] = Array(hi how are you, how is your job, how is your family, how is your brother, how is your sister)

scala> val splitdata = data.flatMap(line => line.split(" "));
splitdata: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[2] at flatMap at <console>:25

scala> splitdata.collect;
res1: Array[String] = Array(hi, how, are, you, how, is, your, job, how, is, your, family, how, is, your, brother, how, is, your, sister)

scala> val mapdata = splitdata.map(word => (word,1));
mapdata: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[3] at map at <console>:25

scala> mapdata.collect;
res2: Array[(String, Int)] = Array((hi,1), (how,1), (are,1), (you,1), (how,1), (is,1), (your,1), (job,1), (how,1), (is,1), (your,1), (family,1), (how,1), (is,1), (your,1), (brother,1), (how,1), (is,1), (your,1), (sister,1))

scala> val reducedata = mapdata.reduceByKey(_+_);
reducedata: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[4] at reduceByKey at <console>:25

scala> reducedata.collect;
res3: Array[(String, Int)] = Array((are,1), (brother,1), (is,4), (sister,1), (family,1), (how,5), (job,1), (you,1), (hi,1), (your,4))

scala>
```

PROGRAM 10

Using RDD and FlMap count how many times each word appears in a file and write out a list of words whose count is strictly greater than 4 using Spark

```
val textFile = sc.textFile("/home/bhoom/Desktop/wc.txt")

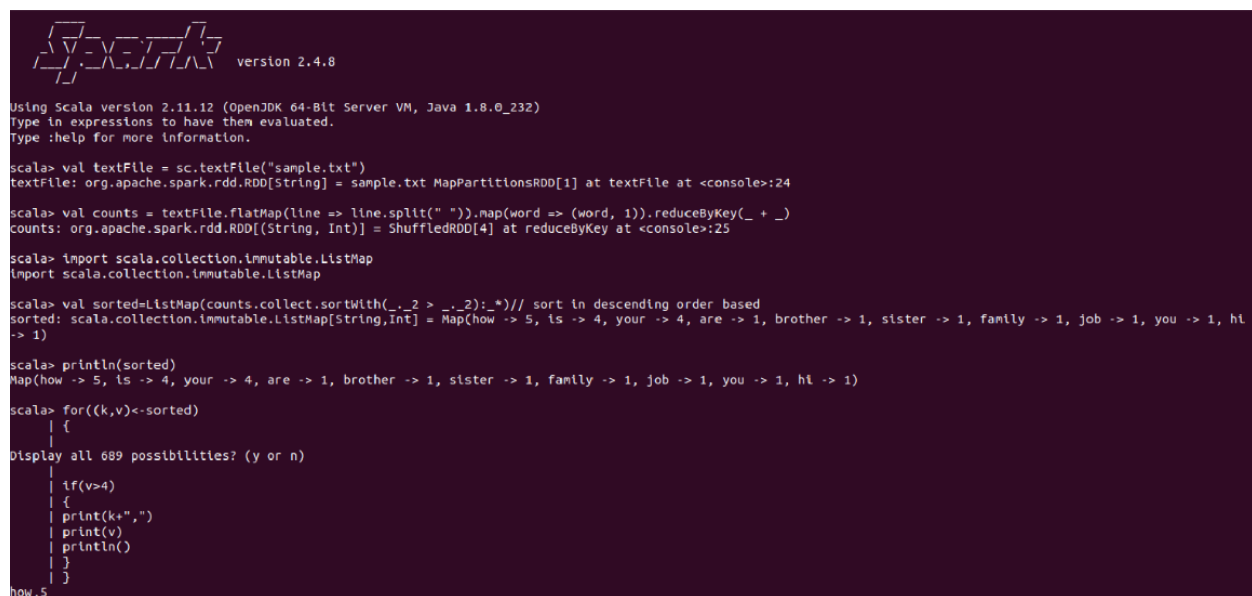
val counts = textFile.flatMap(line => line.split(" ")).map(word => (word, 1)).reduceByKey(_ + _)

import scala.collection.immutable.ListMap

val sorted=ListMap(counts.collect.sortWith(_._2 > _._2):_*)// sort in descending order based
on values

println(sorted)

for((k,v)<-sorted)
{
  if(v>4)
  {
    print(k+",")
    print(v)
    println()
  }
}
```



```
scala version 2.4.8
Using Scala version 2.11.12 (OpenJDK 64-Bit Server VM, Java 1.8.0_232)
Type in expressions to have them evaluated.
Type :help for more information.

scala> val textFile = sc.textFile("sample.txt")
textFile: org.apache.spark.rdd.RDD[String] = sample.txt MapPartitionsRDD[1] at textFile at <console>:24

scala> val counts = textFile.flatMap(line => line.split(" ")).map(word => (word, 1)).reduceByKey(_ + _)
counts: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[4] at reduceByKey at <console>:25

scala> import scala.collection.immutable.ListMap
import scala.collection.immutable.ListMap

scala> val sorted=ListMap(counts.collect.sortWith(_._2 > _._2):_*)// sort in descending order based
sorted: scala.collection.immutable.ListMap[String,Int] = Map(how -> 5, is -> 4, your -> 4, are -> 1, brother -> 1, sister -> 1, family -> 1, job -> 1, you -> 1, hi -> 1)

scala> println(sorted)
Map(how -> 5, is -> 4, your -> 4, are -> 1, brother -> 1, sister -> 1, family -> 1, job -> 1, you -> 1, hi -> 1)

scala> for((k,v)<-sorted)
| {
|   if(v>4)
|   {
|     print(k+",")
|     print(v)
|     println()
|   }
| }
Display all 689 possibilities? (y or n)
how,5
```